

# Introducing Git

Eddie Kovsky  
[ekovsky@redhat.com](mailto:ekovsky@redhat.com)  
IRC: ewk



# An Engineer's Toolbox

- \* Operating system
- \* Programming language
- \* Version control
- \* Your \$EDITOR

THIS IS GIT. IT TRACKS COLLABORATIVE WORK  
ON PROJECTS THROUGH A BEAUTIFUL  
DISTRIBUTED GRAPH THEORY TREE MODEL.

COOL. HOW DO WE USE IT?

NO IDEA. JUST MEMORIZE THESE SHELL  
COMMANDS AND TYPE THEM TO SYNC UP.  
IF YOU GET ERRORS, SAVE YOUR WORK  
ELSEWHERE, DELETE THE PROJECT,  
AND DOWNLOAD A FRESH COPY.



## What changed?

```
$ cat file1.txt
```

```
The Royal Tenenbaums is my favorite movie.  
Luke Wilson is the best Wilson brother.  
He should have won an Oscar for that movie.
```

```
$ cat file2.txt
```

```
The Royal Tenenbaums is my favorite movie.  
Owen Wilson is the best Wilson brother.  
He should have won an Oscar for that movie.
```

# diff

```
$ diff file1.txt file2.txt
2c2
< Luke Wilson is the best Wilson brother.
---
> Owen Wilson is the best Wilson brother.
```

## unified diff

```
$ diff -u file1.txt file2.txt
```

```
--- file1.txt 2019-11-08 11:14:06.927351536 -0700
```

```
+++ file2.txt 2019-11-08 11:14:37.986533112 -0700
```

```
@@ -1,3 +1,3 @@
```

```
The Royal Tenenbaums is my favorite movie.
```

```
-Luke Wilson is the best Wilson brother.
```

```
+Owen Wilson is the best Wilson brother.
```

```
He should have won an Oscar for that movie.
```

## unified diff

```
$ diff -u hello1.py hello2.py
```

```
--- hello1.py    2020-02-10 18:27:16.748772276 -0700
```

```
+++ hello2.py    2020-02-10 18:27:45.261180670 -0700
```

```
@@ -1,4 +1,4 @@
```

```
#!/usr/bin/env python3
```

```
-greeting = "Hello world!"
```

```
+greeting = "Hello gorgeous!"
```

```
print(greeting)
```

# Version control is the history of changes to text.

A version control system (or revision control or source control) is software that records those changes over time.



# An abbreviated list of prior art

- \* Concurrent Versions System – CVS (retired)
- \* Subversion (client-server)
- \* BitKeeper (proprietary)

## How Git is different

```
$ man git
```

```
*NAME*
```

```
git - the stupid content tracker
```

```
*DESCRIPTION*
```

```
Git is a fast, scalable, distributed revision control system
```

# Things we can do with Git

- \* Everyone has a local copy of project history. (**clone**)
- \* Develop features concurrently alongside releases. (**branch** and **rebase**)
- \* Record when a feature was released. (**tag**)
- \* Identify where a bug was introduced. (**bisect**)
- \* Identify who introduced the change. (**blame**)
- \* Quickly search repository history. (**grep** and **log**)

# 2017 kernel rate of change

- \* New kernel released every ~70 days
- \* ~8 changes per hour
- \* 13,000 to 16,000 patches per release
- \* Bonus: Red Hat was the third most active corporate contributor



Provide mechanism,  
not policy.

# Commit Identifier

Hashing the content plus some entropy produces a unique ID.

This is the **Commit Identifier**.

The commit ID is a globally unique number that is mapped to the content change.

commit a430e67d9a2c62a8c7b315b99e74de02018d0a96  
Author: Tadeusz Struk <tadeusz.struk@intel.com>  
Date: Tue Jan 7 14:04:48 2020 -0800

tpm: Handle negative priv->response\_len in tpm\_common\_read()

The priv->response\_length can hold the size of an response or an negative error code, and the tpm\_common\_read() needs to handle both cases correctly. Changed the type of response\_length to signed and accounted for negative value in tpm\_common\_read().

Cc: stable@vger.kernel.org

Fixes: d23d12484307 ("tpm: fix invalid locking in NONBLOCKING mode")

Reported-by: Laura Abbott <labbott@redhat.com>

Signed-off-by: Tadeusz Struk <tadeusz.struk@intel.com>

Reviewed-by: Jarkko Sakkinen <jarkko.sakkinen@linux.intel.com>

Signed-off-by: Jarkko Sakkinen <jarkko.sakkinen@linux.intel.com>

--->8

```
diff --git a/drivers/char/tpm/tpm-dev-common.c b/drivers/char/tpm/tpm-dev-common.c
index b23b0b999232..87f449340202 100644
--- a/drivers/char/tpm/tpm-dev-common.c
+++ b/drivers/char/tpm/tpm-dev-common.c
@@ -130,7 +130,7 @@ ssize_t tpm_common_read(struct file *file, char __user *buf,
                        priv->response_read = true;

                        ret_size = min_t(ssize_t, size, priv->response_length);
-                       if (!ret_size) {
+                       if (ret_size <= 0) {
                                priv->response_length = 0;
                                goto out;
                        }
}
```

```
diff --git a/drivers/char/tpm/tpm-dev.h b/drivers/char/tpm/tpm-dev.h
index 1089fc0bb290..f3742bcc73e3 100644
--- a/drivers/char/tpm/tpm-dev.h
+++ b/drivers/char/tpm/tpm-dev.h
@@ -14,7 +14,7 @@ struct file_priv {
        struct work_struct timeout_work;
        struct work_struct async_work;
        wait_queue_head_t async_wait;
}
```




# Directed Acyclic Graph - The DAG

Commit identifiers are not very useful if we cannot see how they relate to each other.


Git uses a data structure called a graph to build a timeline of commits.

Not this kind of graph




# A graph is a structure with nodes and edges.

Edges connect each node.




Public domain <https://upload.wikimedia.org/wikipedia/commons/thumb/5/5b/6n-graf.svg/333px-6n-graf.svg.png>


In a directed graph, edges have a direction.



An acyclic graph does not have a closed path.



For us humans, time is linear.




# Build and survey history

Nodes are the commit IDs.

Edges connect ancestor commits to construct a timeline.

An (arbitrary) ordering of commits produces a DAG.

# Putting it all together





```

* | | | 9520b835ffda - mfd: 88pm860x-core: Convert to i2c_new_dummy_device (3 months ago) <Wolfram Sang>
* | | | 83215897356f - mfd: 88pm800: Convert to i2c_new_dummy_device (3 months ago) <Wolfram Sang>
* | | | aee36174b22d - mfd: timberdale: Use dev_get_drvdata (3 months ago) <Chuhong Yuan>
* | | | e406b832d89d - mfd: da9063: Remove now unused platform_data (3 months ago) <Wolfram Sang>
* | | | a604e5b29ce6 - mfd: tps80031: Convert to devm_i2c_new_dummy_device (3 months ago) <Wolfram Sang>
* | | | 624e3fceb533 - mfd: cros_ec: Update cros_ec_commands.h (3 months ago) <Yicheng Li>
* | | | ed075453d527 - dt-bindings: mfd: rn5t618: Document optional property system-power-controller (3 months ago) <Jonathan Neuschäfer>
* | | | cbfe612d471f - mfd: aat2870: No need to check return value of debugfs_create_functions (3 months ago) <Greg Kroah-Hartman>
* | | | 64e8a9bacadb - mfd: ab8500: No need to check return value of debugfs_create_functions (3 months ago) <Greg Kroah-Hartman>
* | | | efecc3b531a3 - mfd: ab3100: No need to check return value of debugfs_create_functions (3 months ago) <Greg Kroah-Hartman>
* | | | b5e29aa880be - mfd: davinci_voicecodec: Remove pointless #include (4 months ago) <Arnd Bergmann>
* | | | 197df18f7038 - mfd: max77620: Add of_node_put() before return (4 months ago) <Nishka Dasgupta>
| | | / / / /
* | | | d0b3cfee333e - Merge tag 'backlight-next-5.4' of git://git.kernel.org/pub/scm/linux/kernel/git/lee/backlight (7 weeks ago) <Linus Torvalds>
| | | / / / /
* | | | c0b64faf0fe6 - backlight: pwm_bl: Set scale type for brightness curves specified in the DT (10 weeks ago) <Matthias Kaehlcke>
* | | | 511a204638d7 - backlight: pwm_bl: Set scale type for CIE 1931 curves (10 weeks ago) <Matthias Kaehlcke>
* | | | d55c028f8b25 - backlight: Expose brightness curve type through sysfs (10 weeks ago) <Matthias Kaehlcke>
* | | | 6451e123dec3 - MAINTAINERS: Add entry for stable backlight sysfs ABI documentation (10 weeks ago) <Matthias Kaehlcke>
* | | | ec665b756e6f - backlight: gpio-backlight: Correct initial power state handling (10 weeks ago) <Peter Ujfalusi>
* | | | 28ald72a221e - video: backlight: tosa_lcd: drop check because i2c_unregister_device() is NULL safe (10 weeks ago) <Wolfram Sang>
* | | | bcd69da98e36 - video: backlight: Drop default m for {LCD,BACKLIGHT_CLASS_DEVICE} (10 weeks ago) <Geert Uytterhoeven>
* | | | 30f644fd6ec9 - backlight: lms283gf05: Fix a typo in the description passed to 'devm_gpio_request_one()' (3 months ago) <Christophe JAILLET>
* | | | 072e2c8192cf - backlight: lm3630a: Switch to use fwnode_property_count_uXX() (3 months ago) <Andy Shevchenko>
* | | | 865cac14c2da - backlight: rave-sp: Leave initial state and register with correct device (3 months ago) <Lucas Stach>
| | | / / / /
* | | | 299d14d4c31a - Merge tag 'pci-v5.4-changes' of git://git.kernel.org/pub/scm/linux/kernel/git/helgaas/pci (7 weeks ago) <Linus Torvalds>
| | | / / / /
* | | | c5048a73b477 - Merge branch 'pci/trivial' (7 weeks ago) <Bjorn Helgaas>
| | | / / / /
* | | | 55507aea5882 - PCI: Remove unnecessary returns (2 months ago) <Krzysztof Wilczynski>
* | | | b071c1fd7a8a - PCI: OF: Correct of_irq_parse_pci() documentation (2 months ago) <Lubomir Rintel>
* | | | d1bbf38aaf88 - PCI: Fix typos and whitespace errors (2 months ago) <Bjorn Helgaas>
| | | / / / /
* | | | cdc7e184abb3 - Merge branch 'remotes/lorenzo/pci/vmd' (7 weeks ago) <Bjorn Helgaas>
| | | / / / /
* | | | ala30170138c - PCI: vmd: Fix shadow offsets to reflect spec changes (7 weeks ago) <Jon Derrick>
* | | | e3dffa4f6c36 - PCI: vmd: Fix config addressing when using bus offsets (7 weeks ago) <Jon Derrick>
| | | / / / /
* | | | 3efa7f1febe6 - Merge branch 'lorenzo/pci/tegra' (7 weeks ago) <Bjorn Helgaas>

```

# Let's rewrite history!



## git init

```
$ diff -u file1.txt file2.txt
--- file1.txt 2019-11-08 11:14:06.927351536 -0700
+++ file2.txt 2019-11-08 11:14:37.986533112 -0700
@@ -1,3 +1,3 @@
The Royal Tenenbaums is my favorite movie.
-Luke Wilson is the best Wilson brother.
+Owen Wilson is the best Wilson brother.
He should have won an Oscar for that movie.
```

# Let's rewrite history!

```
$ git rebase
```

```
$ git reflog
```

# The art of finding stuff

Like regular grep, but faster!

```
$ git grep <pattern>
```

```
$ git log --grep=<pattern>
```

```
$ git blame <file>
```

```
$ git bisect
```

# Demo 1

Version 1 is working.

Version 2 has a bug.

Can we find the commit that introduced the bug,  
even if there are hundreds of commits between V1 and V2?

# Let's play a game!

I am thinking of a number between 1 and 100.

**What should your first guess always be?**

# Binary search

If the problem set can be ordered, each pass reduces the problem space by half.



# Demo 1

```
$ git bisect
```

# Demo 2

Customer sees a new error after updating kernel:

```
[drm] Unable to create a private tmpfs mount, hugepage support  
will be disabled(-22).
```

# Demo 2

```
$ git grep '[drm] Unable to create a private tmpfs mount,  
hugepage support will be disabled(-22).'
```

# Demo 2

```
$ git grep '\[drm\] Unable to create a private tmpfs mount,  
hugepage support will be disabled\(-22\)'
```

# Demo 2

```
$ git grep 'Unable to create a private tmpfs mount, hugepage  
support will be disabled'
```

```
drivers/gpu/drm/i915/gem/i915_gem_shmem.c:593:  
DRM_NOTE("Unable to create a private tmpfs mount, hugepage  
support will be disabled(%d).\n"
```

# Demo 2

```
$ git log --grep='Unable to create a private tmpfs mount,  
hugepage support will be disabled'
```

# Demo 2

```
$ git blame drivers/gpu/drm/i915/gem/i915_gem_shmem.c
```

# Demo 3

Find this definition.

```
$ git grep 'struct task_struct'
```



# Demo 3

We only want the definition, but we are getting every usage.

```
$ git grep 'struct task_struct'
```

```
struct task_struct *task;          /* main task structure */  
syscall_get_arch(struct task_struct *task)  
...
```

# Demo 3

Domain specific knowledge of the C language quickly narrows our search:

```
struct foo {  
    field1  
    field2  
    ...  
}
```

# Demo 3

```
$ git grep 'struct task_struct {'
```

```
include/linux/sched.h:629:struct task_struct {  
tools/include/linux/lockdep.h:26:struct task_struct {  
tools/testing/selftests/bpf/progs/strobemeta.h:14:struct task_struct {};  
tools/testing/selftests/bpf/progs/test_core_reloc_kernel.c:24:struct task_struct {
```

# Resources

## man pages

```
$ man git
```

## Pro Git

<https://git-scm.com/book/en/v2>

## Git from the inside out


<https://maryrosecook.com/blog/post/git-from-the-inside-out>

# Thank you


Eddie Kovsky


[ekovsky@redhat.com](mailto:ekovsky@redhat.com)

IRC: ewk

 [linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)

 [facebook.com/redhatinc](https://facebook.com/redhatinc)

 [youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)

 [twitter.com/RedHat](https://twitter.com/RedHat)

